

Semaphores


Daniel Zappala

CS 360 Internet Programming
Brigham Young University


Semaphores

- **semaphore** is a shared variable maintained by OS
 - contains an integer and a queue
 - value initialized ≥ 0
- **wait(s)**: wait for a signal on semaphore *s*
 - decrements semaphore, blocks if value < 0
 - if blocked, process put on the queue, suspends until signal is sent
- **signal(s)**: transmit a signal to semaphore *s*
 - increments semaphore
 - if value ≤ 0 then unblock someone
- **wait()** and **signal()** are atomic operations and cannot be interrupted

wait()



signal()



Types of Sempahores

- **binary semaphore**
 - only one process at a time may be in the critical section
- **counting semaphore**
 - a fixed number of processes > 0 may be in the critical section
- OS determines order that process are released from the queue, but usually FIFO in order to prevent starvation

Using Semaphores

```
1 semaphore s = 1;
2
3 void thread(int i) {
4     while (true) {
5         wait(s);
6         /* critical section */
7         signal(s);
8         /* remainder */
9     }
10 }
```

- semaphore protects critical section
- can set s to > 1 to let more than one process in the critical section
 - $s \geq 0$: number that can enter
 - $s < 0$: number that are waiting

POSIX Semaphores

POSIX Semaphores

```
1 #include <semaphore.h>
2
3 int sem_init(sem_t *sem, int pshared, unsigned int value);
4 int sem_wait(sem_t * sem);
5 int sem_trywait(sem_t * sem);
6 int sem_post(sem_t * sem);
```

- `sem_init()`: sets initial value of semaphore; `pshared = 0` indicates semaphore is local to the process
- `sem_wait()`: suspends process until semaphore is > 0 , then decrements semaphore
- `sem_trywait()`: returns `EAGAIN` if semaphore count is $= 0$
- `sem_post()`: increments semaphore, may cause another thread to wake from `sem_wait()`

Example Code

- see example code `semaphore.cc`

► [GitHub](#)

Producer Consumer

Producer Consumer

```
1 sem_t lock, numItems, numSpaces;
2 sem_init(&lock, 0, 1);
3 sem_init(&numItems, 0, 0);
4 sem_init(&numSpaces, 0, BUFFER_SIZE);
```

producer:

```
1 while (True) {
2     produce();
3     sem_wait(&numSpaces);
4     sem_wait(&lock);
5     append();
6     sem_post(&lock);
7     sem_post(&numItems);
8 }
```

consumer:

```
1 while (True) {
2     sem_wait(&numItems);
3     sem_wait(&lock);
4     take();
5     sem_post(&lock);
6     sem_post(&numSpaces);
7     consume();
8 }
```

Looking at the Code ...

- ① *What is the purpose of semaphore lock?*
- ② *What is the purpose of semaphore numSpaces?*
- ③ *What is the purpose of semaphore numItems?*
- ④ *Why are the semaphores initialized to different values?*
- ⑤ *Can the producer swap the signals for numItems and lock?*
- ⑥ *Can the consumer swap the waits for numItems and lock?*

Important Insights

- two purposes for semaphores
 - **mutual exclusion**: semaphore *lock* controls access to critical section
 - **signalling**: semaphore *numSpaces* coordinates the number of spaces in the buffer, so the producer waits if the buffer is full
 - **signalling**: semaphore *numItems* coordinates the number of items in the buffer, so the consumer waits if the buffer is empty
- avoid race conditions
 - *item* keeps a local copy of the data protected by the semaphore so that it can be accessed later
 - reduces amount of processing inside the critical section
- ordering of semaphores is important