

# Peer-to-Peer Networking

Daniel Zappala


CS 360 Internet Programming  
Brigham Young University

# Definition

- hosts exchange data directly with each other
- hosts act as both clients and servers

# Gnutella

# Napster



- Napster stores a directory of music on your computer, so others can search it, download songs directly from you
- Like sharing cassette tapes or CDs or MP3s with your friends

# Copyright Law

- copyright: owner has exclusive rights to reproduce, adapt, publicly distribute, perform, and display their work
  - **direct infringement**: copying part or all of a copyrighted work without authorization
  - **vicarious liability**: operator has (1) the right and ability to control users and (2) a direct financial benefit from allowing their acts of piracy.
  - **contributory infringement**: requires (1) knowledge of the infringing activity and (2) a material contribution – actual assistance or inducement – to the alleged piracy.

# Fair Use

- use or copying of all or a portion of a copyrighted work without permission of the owner, e.g. for criticism, comment, news reporting, teaching, scholarship, or research
- courts consider:
  - purpose and character of use (commercial vs non profit)
  - nature of work
  - amount and substantiality of portion used (including size and quality)
  - the effect of use on market for or value of copyrighted work
- ▶ A Fair(y) Use Tale

# Napster in Court

- Napster claims they are not infringing copyright because they are not storing any songs
- shutdown by court injunction because case against them was likely to succeed
  - Napster users likely guilty of *direct* copyright infringement - copying of a work by another
  - Napster likely to be guilty of *contributory* infringement because they learned of infringement and failed to purge the materials from its system
  - Napster likely to be guilty of *vicarious* infringement because they supervised or controlled the party engaging in infringing activity and had a financial interest in the activities
- see Wikipedia for background information

# Promotional Power of Free Music


- record companies have claimed that free downloads suppress sales
- some proof of the opposite effect
  - April 2000: tracks from Radiohead's *Kid A* album on Napster three months before CD release
  - millions of downloads by the time the record is released
  - number one spot on the charts in debut week, had never been in the top 20 before
  - beat many other heavily marketed artists
- **this example doesn't excuse piracy, but it does indicate that file sharing can provide a marketing opportunity for new bands**



# Gnutella – version 0.4

- can we share music illegally and not get caught?
- fully distributed, peer-to-peer system
- bootstrapping
  - first time: connect to a peer you heard about outside the system
  - for example, in a chat room
  - keep a cache of all peers discovered and use for bootstrapping next time
- peer discovery
  - try to always be connected to a fixed number of peers (TCP)
  - send a Ping message to existing neighbors, which is flooded to their neighbors
  - other peers respond to Ping with one or more Pong messages, containing IP address, port number, number of files sharing, number of KB sharing

# Gnutella – version 0.4




- queries
  - send a query to your neighbors
  - neighbors flood query, limited by a TTL
  - includes minimum speed in kb/s for responding peers, search criteria
- query hit
  - provide IP address, port, number of hits, speed, result set (file name, size)
  - sent along reverse path

# Gnutella – version 0.4

- download songs directly from peer
- problems
  - no explicit rate limit on ping frequency or query frequency - quickly leads to overload
  - slow peers can hinder faster peers

# Gnutella – version 0.6



- **use hierarchy to scale**
  - super peer: peers with high bandwidth
  - ordinary peer: peers with low bandwidth
- super peers cache names of content held by children
- queries sent among only the super peers


# BitTorrent

# Motivation

- how can an ordinary person, with limited money and bandwidth, serve content to a worldwide audience?
- web servers are limited in their scalability
  - the more clients that need to be served, the slower they access the content
  - eventually the wait becomes so long, TCP connections time out
- solutions
  - Content Delivery Network: spreads the load among a set of servers, but it is expensive
  - Peer-to-Peer File Distribution: spreads the load among a set of peers, inexpensive, must rely on the good will of others

# Modeling File Download

- server upload rate:  $u_s$
- peer upload rate:  $u_i$
- peer download rate:  $d_i$
- file size (bits):  $F$
- total number of peers:  $N$
- assume plentiful bandwidth in the Internet core



# Client-Server Distribution Time

- min download time
  - $\frac{NF}{u_s}$  when constrained by server bandwidth
  - $\frac{F}{d_{min}}$  when constrained by slowest peer,  
 $d_{min} = \min(d_1, d_2, \dots, d_N)$
- $D_{CS} \geq \max(\frac{NF}{u_s}, \frac{F}{d_{min}})$



# Peer-to-Peer Distribution Time

- minimum download time


- $\frac{F}{u_s}$  when constrained by server bandwidth (must deliver the file at least once)

- $\frac{F}{d_{min}}$ , when constrained by the slowest peer

- $\frac{NF}{u_s + \sum_{i=1}^N u_i}$ , when constrained by the overall upload rate

- $D_{P2P} \geq \max\left(\frac{F}{u_s}, \frac{F}{d_{min}}, \frac{NF}{u_s + \sum_{i=1}^N u_i}\right)$


# Comparison



- $F/u_i = 1$  hour,  $u_s = 10u_i$ ,  $d_{min} \geq u_s$
- peer-to-peer download is **self-scaling**: the more peers that download, the more bandwidth is available for upload

# Basic Mechanisms

- 1 download a .torrent file from a web server
- 2 contact the listed *tracker* for a list of peers
- 3 refresh peers as needed
- 4 check with each peer to determine which blocks they have
- 5 parallel download,  $j$  connections, rarest block first



# Incentives

- problem: freeloaders
  - people who try to download without uploading
  - breaks the self-scaling behavior of peer-to-peer distribution
- tit-for-tat
  - serve content to  $k$  connections at a time
  - serve the connections that give you the best download rate
  - periodically serve content to a random connection to see if it can do better than a current connection
  - deny content to all others